

TSM_EmbHardw: Cache and Tightly Coupled Memory

Index

- [1 Aims](#)
 - [2 Lecture Prologue](#)
 - [3 Lecture](#)
 - [4 Lecture Epilogue](#)
 - [5 Exercises](#)
 - [6 Laboratory](#)
 - [Glossary](#)
 - [References](#)
 - [Answers to exercise questions](#)
-

1 Aims

In this lecture we discover some ways to use cache and tightly coupled memory effectively, namely correct alignment, prefetching and tiling for cache as well as some considerations for tightly coupled memory.

2 Lecture Prologue

The concept of cache is an old one. Both [Hardy](#) and [Hennesy and Patterson](#) handle the basics and advanced concepts quite well.

We shall merely cover some aspects.

2.1 Cache - Fundamental Principles

Research in the 60's formulated the principles of **locality** quite quickly. A program spends significant periods of time (**temporal locality**) in sections of code (loops in other words) and data (**spatial locality**).

Putting the following together

- Fetching code/data from RAM is expensive in time
- Principle of locality means that some sections of code/data will be frequently re-used
- Local temporary storage of this code/data will save time and increase performance of the platform
- To preserve performance this local memory should be maintenance free and so transparent to the programmer
- RAM is expensive in space brings us to the idea of cache.

Cache is part of the **memory hierarchy** and typically multilevel L1 -> L3. L1 and L2 are typically in the chip boundaries, L3 external. L1 and L2 are standard on midrange embedded processors.

Cache is commonly understood as a small, hardware managed, memory tightly coupled with the processor holding code and data sparing the processor from avoidable external fetches. To a certain extent cache can be considered a form of **loop optimisation**. The other form of memory tightly attached to the CPU is tightly coupled memory (TCM), or scratchpad memory (SPM). This memory is in the address range of the processor and is not transparent to the programmer.

Cache size and policies affect the hit rate. The **hit rate** is the probability that the looked-for datum is resident in the cache. For example a cache with a hit rate of 80%, a very low number, will execute a read in 2 CPU cycles. For the remaining 20% the CPU must access the system bus at a cost of 3 additional cycles. Therefore, for say 10 instructions the cost is $10 \times 0.8 \times 2 + 10 \times (1-0.8) \times 5 = 16 + 10 \text{ CPU cycles}$. Hence, for nearly 40% of the time the CPU would be unable to get to 20% of the instructions. If the hit rate was to be increased to a more reasonable number such as 90% then the arithmetic would be: $10 \times 0.9 \times 2 + 10 \times (1-0.9) \times 5 = 18 + 5 \text{ CPU cycles}$. The **miss rate** is decreased by half.




Figure 1: Depiction of the basic cache architecture

2.2 Basic Read Operation Cache

Cache operation can be summarised as follows:

- If the CPU requires data or code (data read / instruction fetch) it will generate an address
- If the contents of that address are not in cache the drivers will be activated and external components will be accessed
- This is known as a **cache miss**
- In the case of a cache miss the contents of that address will be placed in cache memory whilst it is being transferred to the CPU
- Address contents are transferred in **cache lines** at the time of writing commonly 64 bytes.
- This is known as a **line-fill**
- The next time the address or a surrounding address, is accessed the contents will already be in cache, representing a **cache hit**
- A **compulsory line fill** is one which could not have been avoided using a better cache policy.

There are four basic operations with respect to cache. A design which allows all signals from the CPU to access the main memory regardless whether it's a hit or a miss is known as a look aside implementation. This makes for faster single-processor read miss accesses since by the time the cache controller realises it has a miss, the system memory is well on its way to getting the required data ready for the CPU.

Look through or in-line caches are typical where the cache controller decides whether to initiate a system memory read/write or not. This kind of implementation massively reduces the CPU-required bandwidth at the cost of the so-called lookup penalty, the time required to decide that a system memory access is required. In a read hit the system bus address and data buffers remain switched off and the Cache Data Memory buffers are switched on, allowing the datum to proceed to the CPU as if it were coming from system memory, only faster. During a read miss the opposite is true – after the lookup penalty the system bus address and data buffers are turned on and those of the Cache Data Memory to write.




Figure 2: Depiction of the principles of reading operations with cache

2.3 Basic Write Operation Cache

Code or read-only data that is in cache can be discarded when it becomes out of context. Written data necessitates transfer to in main memory once the cache line is to be otherwise occupied. Write hits are generally handled in two ways:

Write through -uUpdate of both the Cache Data Memory and the System memory avoiding the issue of cache coherence. Cache coherence is the issue when there are two different versions of the datum at any one period of time in the system. This is known as a write-through policy.

Copy back policy is to only write the datum into the cache – speed at the expense of management effort to ensure cache coherence. The simplest method – and there are lots of others - is to implement a dirty bit. The dirty bit signifies that a cache line has been written to so that when the time comes for the cache line to be replaced resident data is **evicted**.

Write misses are also handled in a variety of ways. For instance, and not complete by far:

- The cached line can be ignored and the result written directly to memory,
- The cached line is written to both memory and cache
- The line is simply written to cache and another policy implemented after that
- Update the cache line regardless of whether there is a cache hit or miss.




Figure 3: Depiction of three principles of write operations with cache

2.4 Cache Bypassing for I/O

There is also a strong case to be made for ignoring cache entirely and implementing an uncached addresses. This is useful in systems where I/O is memory mapped.




Figure 4: Depiction of the principles of I/O operations with cache

3 Lecture

[Back To Index](#)

Having recalled the basics of cache operation we now need to look at how to use it properly.

We shall look at

1. Alignment
2. Prefetching
3. Tiling
4. Cache locking

Then we shall discuss 5. Tightly Coupled Memory

as well as some relevant case studies

3.1 Alignment

(This section intentionally left blank)

3.2 Notes to Cache Prefetching

(This section intentionally left blank)

3.3 Notes to Tiling

(This section intentionally left blank)

3.4 Notes to Cache Locking

(This section intentionally left blank)

3.5 Tightly Coupled Memory

See for instance [Panda 2002](#) & [Panda 2000](#)

4 Lecture Epilogue

[Back To Index](#)

4.0 Cache Generalities

Some processors support instructions for cache manipulations. The GNU compiler [GNU-3](#) offers a `--builtin` function to clear the code cache Linux [MAN-1](#) offers a function to flush either cache

```
#include <sys/cachectl.h>

int cacheflush(int nbytes;
               void addr[nbytes], int nbytes, int cache);
```

4.1 Cache Alignment

The alignment problem is twofold

- The arrays (input data) are not aligned on cache-line boundaries
- Two arrays that may be used in a calculation cause cache use conflicts

Real-World Support The language AND/OR compiler can be given hints on preferred placement. for instance

```
#include <stdalign.h>

// C11-> C23
_Alignas(expression)

// C23 ->
alignas(expression)
```

The compiler generally determines allocation in the address space and interference with this principle is not generally preferred.

```
int x __attribute__((aligned (16))) = 0;
```

See: [GNU-1](#)

This is however an alignment to a word size (affecting f.i. bus accesses), not explicitly to a cache line boundary which is a *runtime* attribute. If you know the cache size - most ways of finding this out are runtime specific - then you may be willing to take the chance.

With two arrays it is the relative positioning that is important. Both arrays will have different tags but if the indexes are the same (and both are cache aligned) then they will overlap, and that can be problematic. The ways of alleviating this are:

- if the size of the individual arrays are smaller than the cache then making a two-dimensional array or a structure with two arrays can be the way to go
- if the size of the individual arrays are larger than the cache then the introduction of padding may alleviate the issue

For further thoughts see: [Hong](#)

4.2 Cache Pre-Fetching

Real-World Support

The GNU compiler offers a built-in function (as previously seen in custom instructions)

```
void __builtin_prefetch (const void *addr, ...);
```

See [GNU-2](#) The GNU compiler also offers prefetching optimisations: See [GNU-3](#)

Useful Literature The article from [Vanderwiel](#) is well worth consulting especially for its ideas on hardware pre-fetching. For further thoughts see: [Kuehn](#)

4.3 Tiling

There is an exhaustive body of work for varying application domains on this theme - IEEExplore is a useful starting point.

In the context of dedicated GPU architectures a **tiler** was often included in the GPU architecture to ensure that the data was efficiently collected into the GPU and processed there.




Figure 5: Tiling in an ARM Utgard GPU

From: [ARM](#)

4.4 Cache Locking

The operation of locking the cache is very useful in n-way associative caches since several routines (code) or arrays (data) can be locked in cache and yet enough ways are available for "random" cache accesses. This helps avoid manipulating the run-time memory map at compile-time. The sequence for achieving this – typical for any locking sequence is: to

1. Invalidate the caches
2. Load the caches in some way
3. Lock the cache

For an imaginative and entertaining way of illustrating cache locking we can consider the MPC8349e. The MPC8349e implements data and instruction caches of 32kBytes each with a line size of 8 words and are both 8-way associative. Loading and locking the cache is achieved as follows:

In the case of instructions these are done using a speculative fetch. For this microcontroller, an instruction sequence of a divide followed by a branch is carried out. The divide because it needs lots of cycles during which the code at the destination of the branch is loaded speculatively into cache. If the branch is not fulfilled then whilst these instructions will not be executed, they are in cache and marked valid. It is then possible to lock either the entire instruction cache or simply this specific way.

4.5 Tightly Coupled Memory

(This section intentionally left blank)

5 Exercises

[Back To Index](#)

Exercise 1 - Associative Cache Addressing

Assume a 4-way cache of block size 32 bytes and total size 8 kByte on a 32-bit architecture. What are the number of offset, set and tag bits?

[Answer](#)

Exercise 2 - Cache Terminology

Given the following code:

```
#define MAXSIZE 32

void swap (int *, int *, int);

#pragma location = 0x10000000
int a[MAXSIZE] = {1, 2, 3, ... 30, 31, 32};

function (int *a) {
    int temp;
    temp = a[0];      // line A
    return temp;
}
```

How often is the cache flushed?

[Answer](#)

Exercise 3 - Cache Hit and Miss Rates

A 100MHz. processor system has a two layer cache system with the following characteristics

Cache Layer 1 access time $T_{c1} = 4$ clock cycles Cache Layer 2 access time $T_{c2} = 6$ clock cycles Main memory access time $T_m = 8$ clock cycles

Assume a hit rate of 90% for cache layers 1 and 2. What's the average access time?

Hint: use the equations

$$T_{a2} = T_{c2} + (1-h) * T_m$$

$$T_{a1} = T_{c1} + (1-h) * T_{a2}$$

which give you the average access time of the second and first layer cache resp.

[Answer](#)

Exercise 4 - Prefetching

The processor has a direct mapped data cache with lines of 32 bytes and an integer size of 32 bits. Given the following code:

```
#define MAXSIZE 32

void swap (int *, int *, int);

#pragma location = 0x10000000
int a[MAXSIZE] = {1, 2, 3, ... 30, 31, 32};

#pragma location = 0x20000010
int b[MAXSIZE] = {32, 31, 30, ... 3, 2, 1};

swap (&a, &b, MAXSIZE) {
    int temp;
    for (int i = 0; i<=MAXSIZE; i++) {
        temp = a[i]; //line C
        a[i] = b[i]; //line D
        b[i] = temp; //Line E
    }
}
```

1. How many compulsory cache misses does this code generate?
2. Re-locate the arrays to give each a unique location in cache whilst retaining the tag address.

3. Use prefetches to eliminate the compulsory cache misses – discuss qualitatively how much faster the solution is. Explain your solution in terms of useless pre-fetches and cache

[Answer](#)

6 Laboratory

(This section intentionally left blank)

[Back To Index](#)

Glossary

(This section intentionally left blank)

References

Hardy J. Hardy "The Cache memory book". Second Edition, Academic Press, San Diego 1998.

Hennessy and Patterson 2.5 Cross-Cutting Issues: The Design of Memory Hierarchies. Hennessy, John L.; Patterson, David A.. Computer Architecture: A Quantitative Approach (The Morgan Kaufmann Series in Computer Architecture and Design) (p. 292). Elsevier Science. Kindle Edition.

GNU-1 'Specifying Attributes of Variables' <https://gcc.gnu.org/onlinedocs/gcc-3.2/gcc/Variable-Attributes.html> Last accessed 10.10.2025

GNU-2 '7.12 Other Built-in Functions Provided by GCC' <https://gcc.gnu.org/onlinedocs/gcc/Other-Builtins.html> Last accessed 10.10.2025

GNU-3 '3.12 Options That Control Optimization' <https://gcc.gnu.org/onlinedocs/gcc/Optimize-Options.html> Last accessed 10.10.2025

Hong Changwan Hong, Wenlei Bao, Albert Cohen, Sriram Krishnamoorthy, Louis-Noël Pouchet, Fabrice Rastello, J. Ramanujam, and P. Sadayappan. 2016. Effective padding of multidimensional arrays to avoid cache conflict misses. In Proceedings of the 37th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI '16). Association for Computing Machinery, New York, NY, USA, 129–144. <https://doi.org/10.1145/2908080.2908123>

Kuehn-3 Roland Kühn, Jan Mühlig, and Jens Teubner. 2024. How to Be Fast and Not Furious: Looking Under the Hood of CPU Cache Prefetching. In Proceedings of the 20th International Workshop on Data Management on New Hardware (DaMoN '24). Association for Computing Machinery, New York, NY, USA, Article 9, 1–10. <https://doi.org/10.1145/3662010.3663451>

Stock Stock, Gregory and Hahn, Sebastian and Reineke, (2019) 'Cache Persistence Analysis: Finally Exact' Proceedings of the 2019 IEEE Real-Time Systems Symposium (RTSS), <http://dx.doi.org/10.1109/RTSS46320.2019.00049>, <https://arxiv.org/abs/1909.04374>

ARM-1 'Utgard Shader Core' <https://developer.arm.com/documentation/102538/0100/Utgard-Shader-Core> Last accessed 20.10.2025

Vanderwiel Steven P. Vanderwiel and David J. Lilja. 2000. Data prefetch mechanisms. ACM Comput. Surv. 32, 2 (June 2000), 174–199. <https://doi.org/10.1145/358923.358939> <http://www.ece.lsu.edu/tca/papers/vanderwiel-00.pdf>

Panda 2000 Preeti Ranjan Panda, Nikil D. Dutt, and Alexandru Nicolau. 2000. On-chip vs. off-chip memory: the data partitioning problem in embedded processor-based systems. ACM Trans. Des. Autom. Electron. Syst. 5, 3 (July 2000), 682–704. <https://doi.org/10.1145/348019.348570>

Panda 2002 Preeti Ranjan Panda and Nikil D. Dutt. 2002. Memory Architectures for Embedded Systems-On-Chip. In Proceedings of the 9th International Conference on High Performance Computing (HiPC '02). Springer-Verlag, Berlin, Heidelberg, 647–662. (preprint) <https://ics.uci.edu/~dutt/pubs/bc12-hipc02-panda.pdf>

Answers to exercise questions

[Back To Index](#)

Answer to exercise 1 - Associative Caches

The addressing scheme is as follows:

Tag	Set	Offset
-----	-----	--------

Offset: 32 bytes -> offset of $2^5 = 5$ bits

Set: 4 ways means each address represents $4 * 32 \text{ bytes} = 128 \text{ bytes}$

Cache is 8 kByte large, over four ways in blocks of 128 Bytes -> $8 \text{ kB} / 128 \text{ B} = 64$ thus $2^6 = 6$ bits

Tag: the rest of the higher order address bits make up the tag address -> $32 - 5 - 6 = 21$.

Answer to exercise 2 - Cache Terminology

In the context of the function never as the contents of cache are not overwritten In the context of the loading of the array contents, if this overwrites a line that has been made dirty from another part of the program or process then that line is flushed before being overwritten.

Answer to exercise 3 - Hit and Miss Rates

$$T_{a2} = 60 \text{ ns} + (0.1) * 8 * 10 \text{ ns} \rightarrow 68 \text{ ns}$$

$$T_{a1} = 40 \text{ ns} + (0.1) * 68 \text{ ns} \rightarrow 46.8 \text{ ns}$$

Answer to exercise 4 - Prefetching

1. Line size is 32 bytes = 8 ints -> array requires 4 lines of cache -> 8 compulsory cache misses for the two arrays and one for the variable temp.
- 2.

```
#pragma location = 0x10000000
int a[MAXSIZE] = {1, 2, 3, ... 30, 31, 32};
#pragma location = 0x20000080
int b[MAXSIZE] = {32, 31, 30, ... 3, 2, 1};
```

3.

```
#define MAXSIZE 32

void swap (int *, int *, int);

#pragma location = 0x10000000
int a[MAXSIZE] = {1, 2, 3, ... 30, 31, 32};

#pragma location = 0x20000080
int b[MAXSIZE] = {32, 31, 30, ... 3, 2, 1};

swap (&a, &b, MAXSIZE) {
    int temp;
    prefetch(temp);
    prefetch(a[0]);
    prefetch(b[0]);

    for (int i = 0; i < MAXSIZE; i+=8) {
        prefetch(a[i+8]);
        prefetch(b[i+8]);

        temp = a[i]; // unroll 8 times
        a[i] = b[i]; //
        b[i] = temp; //

        temp = a[i++]; //
        a[i] = b[i]; //
        b[i] = temp; //

        .....
    }
}
```

The loop is obviously faster because its overhead is only executed four times instead of 32 times due to loop unrolling. The prefetches should see to it that there are no compulsory misses. One could argue that the prefetches might be better placed near the end of the loop (according to the optimal pre-fetch distance) to reduce the chances of cache eviction occurring and so, the prefetch being rendered useless. The prefetches executed in the final iteration are clear cases of cache pollution.